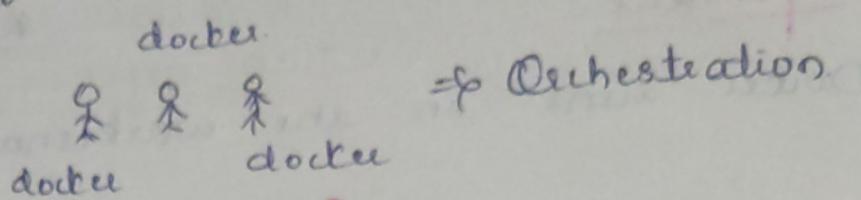


KUBERNETES

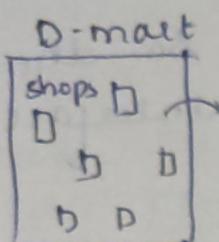
Kubanekes

- Google websites was crashing, lot of traffics do autoscale & autoheal kubernetes came into existence
 - Ex: phie hea phei.



* Why do we use kubernetes?

Monolithic



contains
stored in
one repository

- If anything crashes the entire system should be updated

- Huge cost, management problem

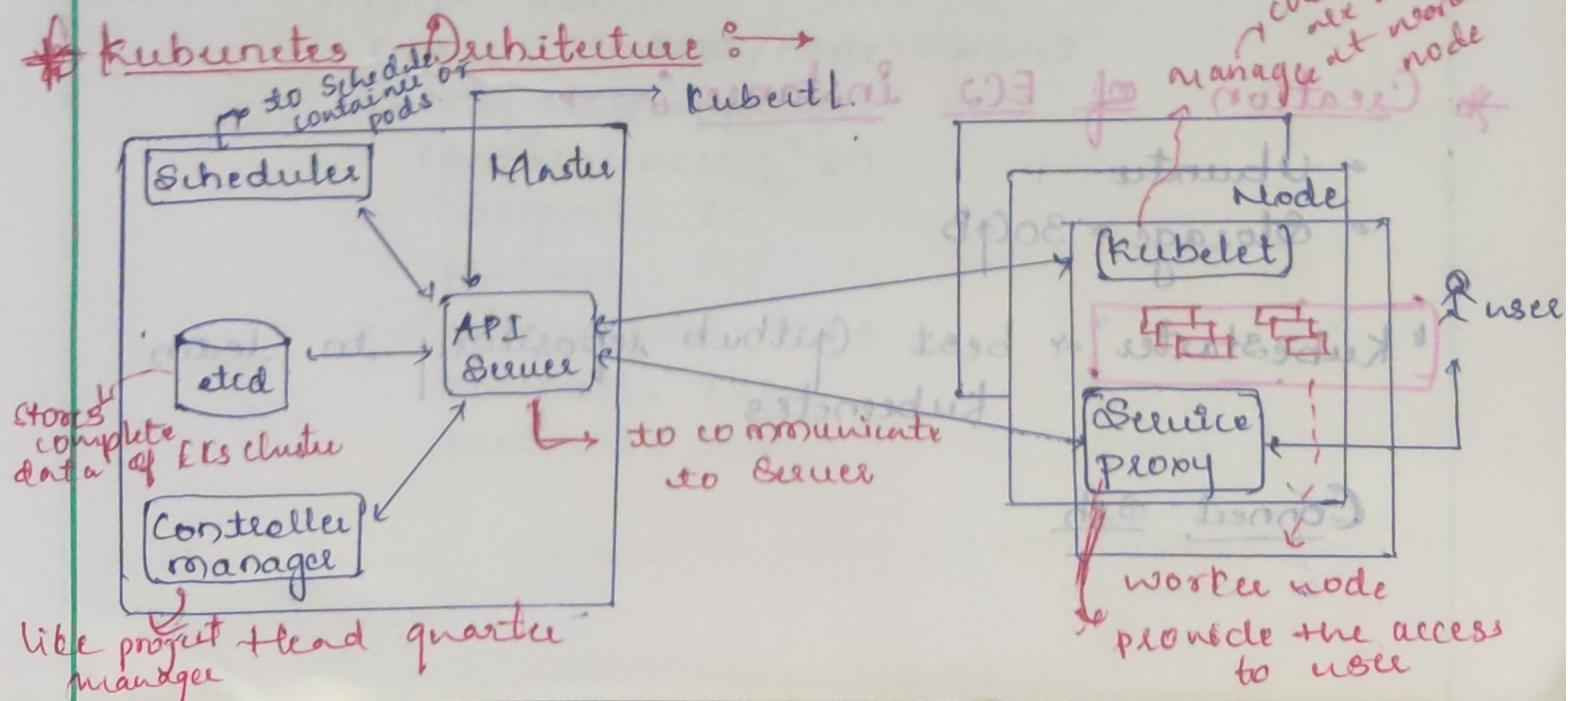
- + Many companies use microservices + Kubernetes (dev and ops)

Inode = Server

Multinode- Cluster (Multiple Server)

↳ Kubernetes handle \uparrow $\text{ip} \backslash \text{exit} \backslash \text{exit}$

~~# kubernetes Architecture :-~~



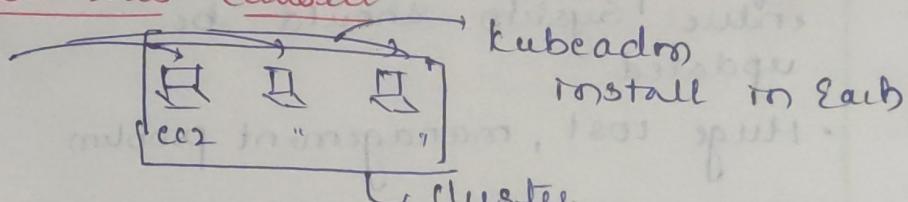
- Cognizant main branch is US - Headquarter (Master)
 - Sub branches pune, bangalore, hyderabad, Noida

Important

- Docker runs on the worker node → Masters make them to work
- API Server communicate with the worker node
- Server cannot communicate directly so the API is used (Application programmable interface) to communicate.

Nodes

- * Creation of Kubernetes Cluster
 - Kubeadm
 - install in each node in "minion"
 - cluster (Heavy & more cost)
 - Minitube (local/ec2)
 - kEND cluster
 - kubernetes in docker
 - EKS/AKS/GKE



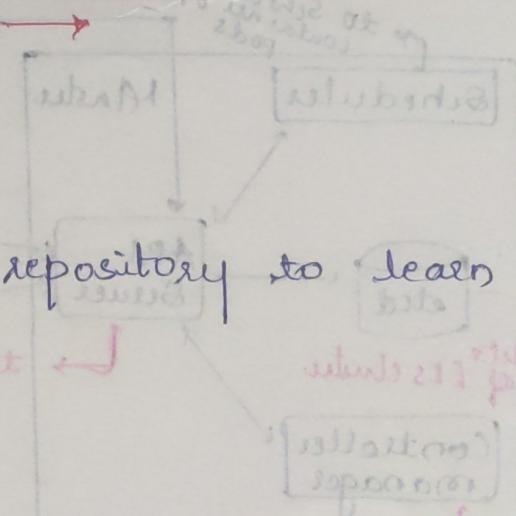
Creation of EC2 Instance

- Ubuntu
- Storage = 30GB

"Kubestarter" → best Github repository to learn Kubernetes

Connect ssh

show ec2now
region with availability zones at



aws ec2 start-instance-id operation

* Kind cluster for Macos & Linux
After installation of docker
docker ps = user denied
To give permission
Command = `Sudo usermod -aG docker $USER && newgrp docker.`

* Minikube for windows installation.

ubuntu = install kind

kubectl

docker

* mkdive kind-cluster

cd kind-cluster

vim config.yaml

key value
kind: cluster

apiVersion: kind.x-k8s.io/v1alpha4

nodes:

- role: control-plane
list image: kindest/node: v1.31.2

- role: worker
image: kindest/node: v1.31.2

- role: worker
image: kindest/node: v1.31.2

- role: worker
image: kindest/node: v1.31.2

extraPortMappings:

- containerport: 80
hostport: 80

protocol: TCP

- containerport: 443
hostport: 443

protocol: TCP

* kind Create cluster --name tuas-cluster --config = config.yaml.

* Minikube Installation Guide for Ubuntu

1. update system packages

```
Sudo apt update
```

2. Install required packages

```
Sudo apt install -y curl wget apt-transport-  
-https
```

3. Install docker

```
Sudo apt install -y docker.io
```

Start & Enable docker

```
Sudo systemctl enable --now docker
```

Add current user to docker group

```
Sudo usermod -aG docker $USER && newgrp  
docker
```

4. Install minikube

```
curl -LO minikube https://storage.googleapis.  
com/minikube/releases/latest/minikube-linux-  
amd64
```

```
chmod +x minikube
```

```
Sudo mv minikube /usr/local/bin/
```

5. Install kubectl

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s  
https://dl.k8s.io/release/stable.txt)/bin/linux/  
amd64/kubectl"
```

Check above image executable & move it
into your path

```
chmod +x kubectl
```

```
Sudo mv kubectl /usr/local/bin/
```

6. Start Minikube

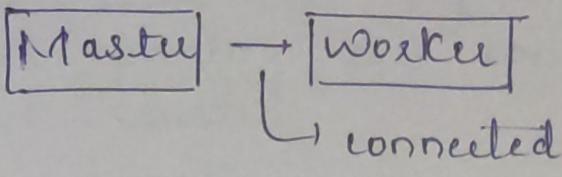
```
minikube start --driver=docker --vm=true
```

7. check cluster status

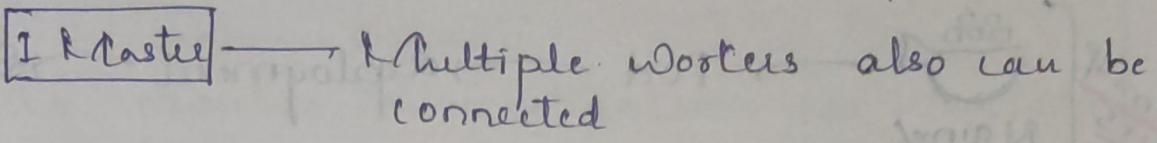
```
minikube status
```

→ uses container in the background to create Clusters

* Kubeadm :



6443-port



kubeadm init

↳ it becomes master

Create EC2 instances

One Master

connect ssh

2 terminals master & worker

install by using the commands = kubestarter

* Both uses contained it creates the requirements of each

- Master

etcd

Scheduler

API

Manager

- Worker

same

- Worker

nodes

Kubelet

Service proxy

gets installed

- If in this whatever I do kubeadm init
it becomes master

* Kubeadm token create --print-join-command

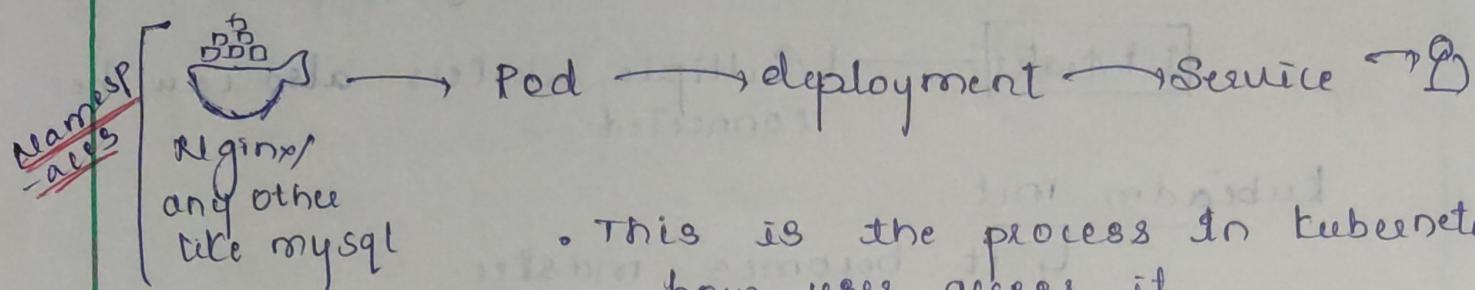
↳ this command removes kubeadm , this is beneficial for worker node

↳ no need of kubeadm

- If in case I did "kubeadm init" at worker mode this command helps to recover that mistake.

+ kubectl get nodes

Peds :-



- This is the process in Kubernetes how user access it

↳ Same as groups \Rightarrow they are isolated
the upper pod do not disturb to another pod.

~~mkd~~ mkdie kubernetes

cd " today

minitube Start --opus 2

- Minitube installation
 - wsl config

kubectl get namespace

$$pod = \text{qginx}$$

kubectl create ns argos -l app=argos

Kubernetes get ns

```
kubectl run nginx --image=nginx
```

kubectl get pods -t --selector=task=knobbe

kubectl delete nginx -ns

* Kubernetes pods → pods are Kubernetes objects that are the basic unit for running our containers inside our Kubernetes cluster.

* Kubernetes Namespaces :-

namespaces provide a mechanism for isolating groups of resources within a single cluster.

• mk dir nginx

• cd "

• vim nginx.yaml
namespace:

yaml = manifest file

kind : Namespace

apiVersion: v1

metadata:

kubectl apply -f namespace.yaml

vim pod.yaml

kind: pod

apiVersion: v1

metadata:

name: nginx-pod

namespace: nginx

Spec:

containers:

- name: nginx

image: nginx:latest

ports:

- containerPort: 80

kubectl apply -f pod.yaml

kubectl get pods -n nginx

kubectl exec -it pod/nginx-pod -n nginx -- bash

- + To ensure that the pod is running correctly or not

Kubectl describe pod/nginx-pod -n nginx

↳ provides complete information

Kubectl exec -it nginx-pod -n nginx -- bash
exit

- Vim deployment.yaml => Search on chrome deployment kubernetes

kind: Deployment

apiVersion: apps/v1

metadata:

name: nginx-deployment

namespace: nginx

Spec:

replicas: 2

selector:

matchLabels:
app: nginx

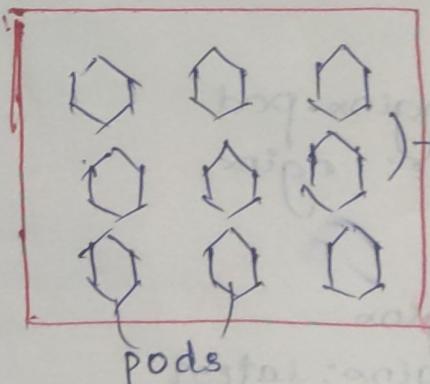
template:

metadata:

labels:

app: nginx

Replicaset / StatefulSet / Deployment



All 3 maintains replica b.

Eg: On Big billion days Amazon gets lot of traffic & the Amazon has only 3-pods now. It will create the replicas to store control traffic.

Only 1 or 2 pods cannot handle the traffic so the replicas are created.

• pod contains replication controller which controls the replicas of pods

* **Replica Set** - Kubernetes resource that ensures a specified number of identical pod replicas are running at any given time.

Ex I created a template

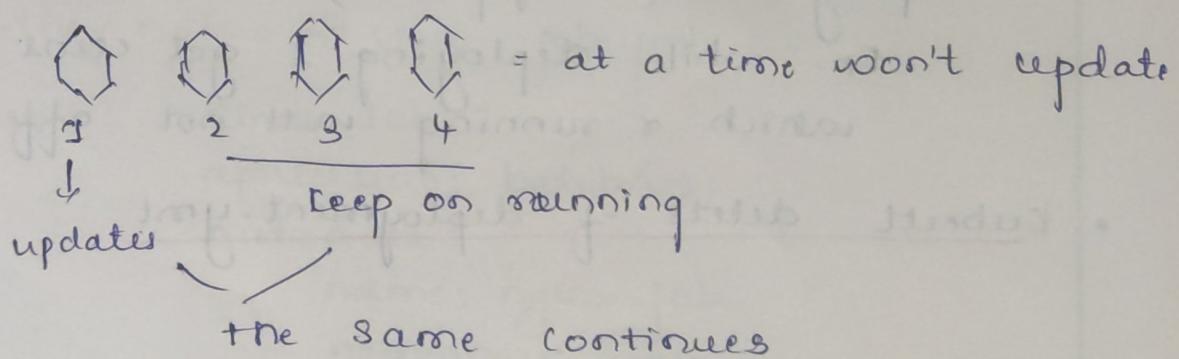
Specified $\text{replicas} = 4 \Rightarrow$ it creates 4 replicas

* **Statefulset** - It mainly keeps the record of pods like running or deleted. It provides the numbers to each pod to maintain the sequence.

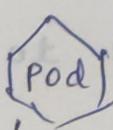
* **Deployment** - It mainly handles the rollback or rolling updates

Ex. The image present in pods, I changed them. It gets changed to all other pods by due to replica.

- In this process we get downtime, System doesn't work.



* **Labels & Selectors** →



I named it as

nginx
label.

match & manage specific pods based on their labels



match the labels

→ no need bcz I have mentioned
its works in deployment.yaml

- kubectl get pods -n nginx
- kubectl delete -f pod.yaml
- kubectl apply -f deployment.yaml
- kubectl get deployment -n nginx
 - " " pods -n "
- kubectl scale deployment/nginx-deployment -n nginx --replicas=5
 - ↓
 - 5 will be created
 - '0' number ?
- kubectl get pods -n nginx
- kubectl set image deployment/nginx-deployment -n nginx=nginx:1.26.3

Deployment (rollingupdate) ↗

Ex- while deploying I get error, the pods
which are running will not affect

- kubectl delete -f deployment.yaml

* cp deployment.yaml replicaset.yaml

↳ copy the same yaml file to replica file.

ls

view replicaset.yaml

↳ the same file is created as deployment.

kubectl apply -f replicaset.yaml

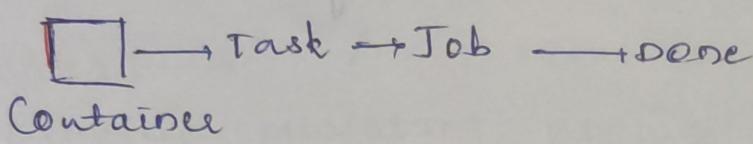
kubectl get replicaset -n nginx

* Daemonsets → relica or copy of

Ensures that atleast 1 pod should be running on node atleast each node

- vim daemonsets.yaml
- cp replicaset.yaml daemonsets.yaml
- kubectl apply -f daemonsets.yaml
- kubectl get pods -n nginx
- kubectl delete -f daemonsets.yaml

* Jobs →



Ensures that a pod runs to completion successfully

- Ex: update shutdown
- Only once run & stop
- It's a parallel process
- 8-pods running parallelly

/kubernetes/nginx → vim job.yaml

kind: Job

apiVersion: batch/v1

metadata:

name: nginx-job

namespace: nginx

Spec:

completions: 1

parallelism: 1

template:

metadata:

name: demo-job-pod

labels:

app: batch-task

Spec:

containers:

• name: batch-containee

image: busybox: latest

command: ["sh", "-c", "echo Hello, Posto!"]

& sleep 30"]

restartPolicy: Never

* **kubectl apply -f job.yaml**

* kubectl get job -n nginx

" " pods " "

running → Created → Stop

If you want back running = delete & start

* **Cron job:**

Following particular Schedule

- Is used to run jobs on a scheduled basis

View cronjob.yaml

kind: CronJob

spec.cronSchedule: batch/v1

metadata:

name: minute-backup

namespace: nginx

Spec:

Schedule: "*/* * * * *

JobTemplate:

Spec:

Template:

Metadata:

name: minute-backup

Labels:

app: minute-backup

* * * * *
min hr day day month
hour week

Spec:

contains:

- name: backup - container

- image: busybox

command:

- sh

- -c

- >

```
echo "Backup started";
```

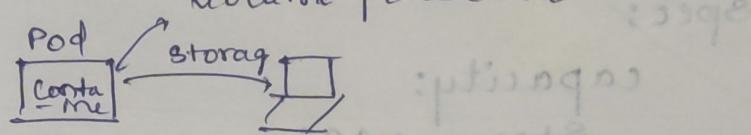
```
mkdir -p/backups &&
```

```
mkdir -p/data &&
```

```
cp -r /demo-data/backups
```

* Volumes →

provide persistent storage to containers inside pods

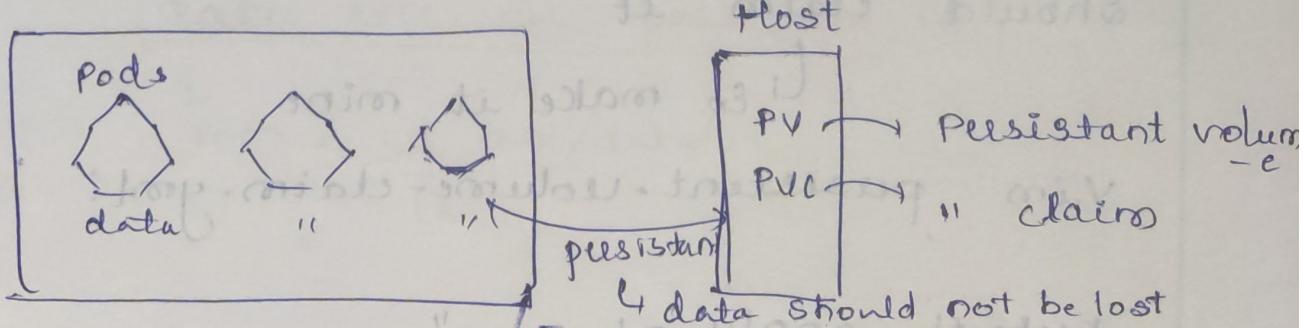


for Cronjob

```
kubectl get -n nginx pods -- Every sec  
" Enter
```

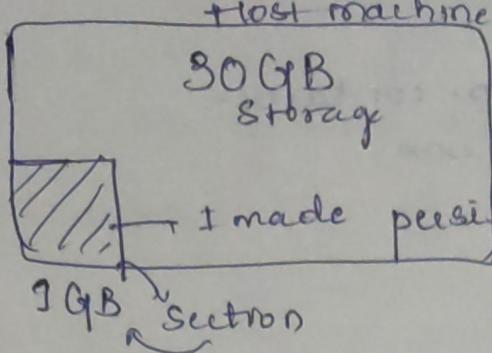
```
kubectl delete -f cron-job
```

* Storage →



- The pods running on nodes if it gets crashed or deleted so there should be persistent storage → ↗

Persistent Volume



80 GB
Storage

3 GB
section

I made persistent volume of 1GB from host machine

```
vim persistent-volume.yaml
```

kind: PersistentVolume

apiVersion: v1

metadata:

name: local-pv

labels:

app: local

spec:

capacity:

storage: 1Gi

accessModes:

- ReadWriteOnce

persistentVolumeClaimPolicy: Retain

storageClassName

```
* kubectl apply -f persistent-volume.yaml
```

* Persistent Volume Claim →

Now I have persistent volume but I should claim it

↳ Ex: make it mine

```
vim persistent-volume-claim.yaml
```

kubectl apply -f "

kubectl get pvc

contains

VolumeMounts:

- mountPath: /var/www/html
name: local-pvc

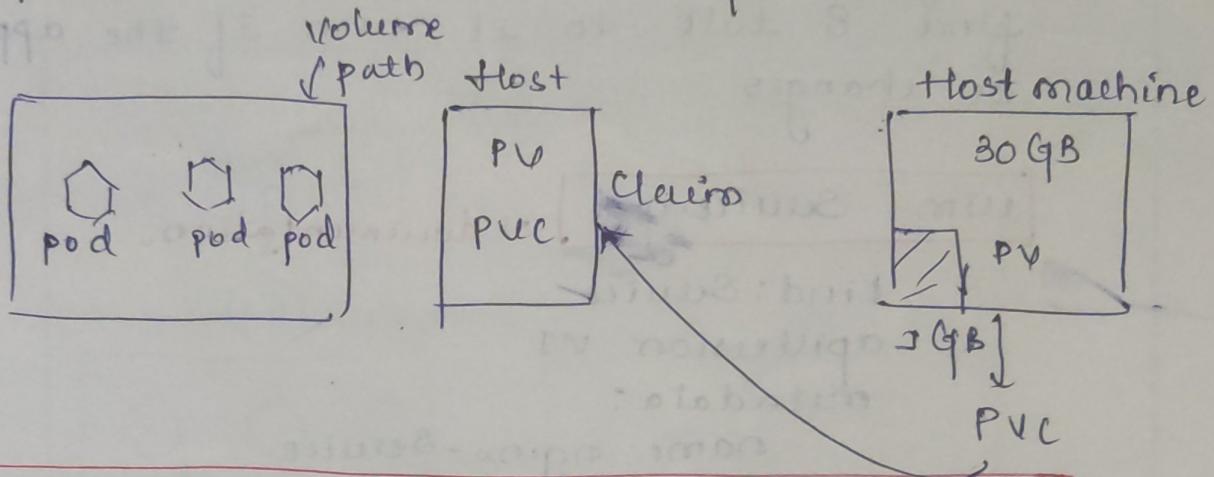
Volumes:

- name: local-pv

PersistentVolumeClaim:

ClaimName: local-pvc.

- Local-path is bound with persistent-volume



- Kubectl get -f persistent-nginx -o yaml

↓
by this we can check
the pod is working on
which node.

- As I'm connected to EC2 the worker node is inside the docker

docker ps

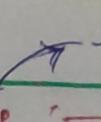
↳ search for worker node

docker exec -it paste the id.

↳ root: cd mnt/data/

ls

↳ few files because bootstrap will get the
[transaction store will take 3-4 mb]

 helps to connect pods reliably
Even if pods come & go, the Service always

* Service: → gives you a fixed way to reach them.

↓ redirect through deployments

 A Service is a method for exposing a network application that is running as one or more pods in your cluster

Ex. It is like a helper that gives your app(pod) a fixed name & address so others can always find & talk to it, even if the app restarts or changes.

File Service.yaml → documentation.

kind: Service

apiVersion: v1

metadata:

name: nginx-Service

namespace: nginx

Spec:

selector:

app: nginx

port:

- protocol: TCP

port: 80

targetPort: 80

kubectl apply -f Service.yaml

kubectl get all

[bc it is in docker]

→ it can't access by EC2

kubectl port-forward service/nginx-Service

-n nginx 80:80 --address=0.0.0.0

Sudo -E "

↳ if this command doesn't work write
Sudo -E & it will create environment

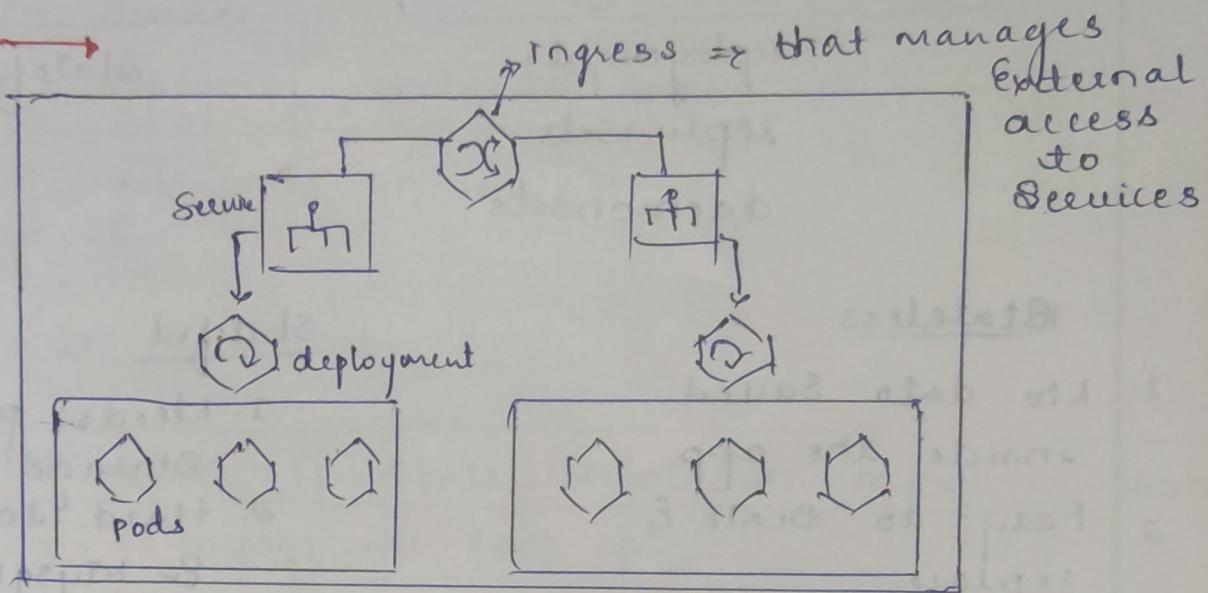
cd ..

Access the IP address on EC2
[git clone "paste"]
cd django-notes ↳ mini-project
ls ↳ git check-out dev

docker build -t notes-app-l8s

* docker-hub

* Ingress :



• Using ingress.yaml

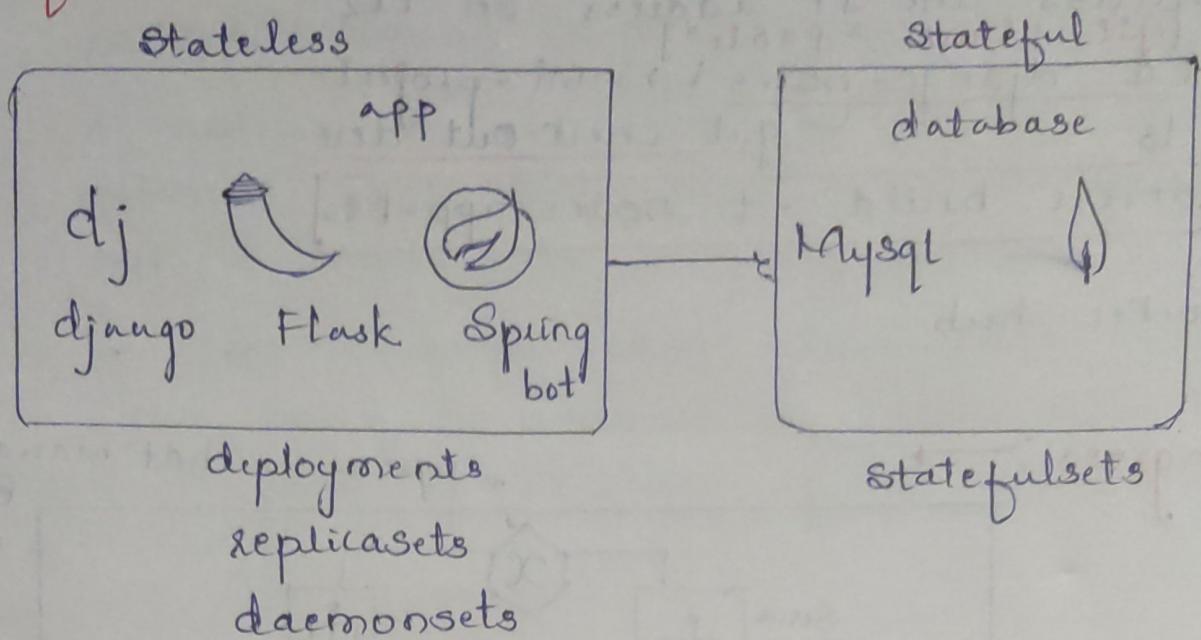
* Annotations :

- They are like sticky notes attached to your Kubernetes objects.
- They store extra information that tools or humans need, but Kubernetes itself doesn't use them to make decisions.

* Cluster Networking :

All the parts of your Kubernetes System talk to each other - like pods, services & nodes

* Statefulsets:



Stateless

1. No data saved inside the app
2. Easy to scale & replace
 - Ex: web server, API gateway

Stateful

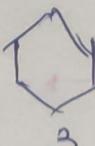
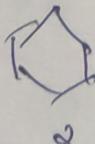
1. Needs persistent storage
2. Hard to scale
 - Ex: MySQL, MongoDB

* Statefulsets ↪

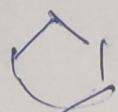
Controllers used to manage applications that require a stable, unique identity for each pod & persistent storage

Ex: I have 3-pods

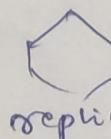
Benefit



if this deletes it will
create replica which will
be the same as before



= data deleted



the same data will
replica get replicated.

cd kubernetes / cd mysql =

kubernetes / Mysql = vim Statefulsets.yaml (documentation)

Kind: Statefulset

application: apps/v1

metadata:

name: mysql-statefulset

namespace: mysql

Vim namespace.yaml

Vim service.yaml

• kubectl apply -f statefulset.yaml

• kubectl get pods -n mysql

• watch "

↳ 0 Statefulset
1
2

• kubectl exec -it statefulset-0 -n mysql -- bash

@root@mysql ~

> mysql = Show database

• kubectl delete pod mysql-statefulset-0 -n mysql

↳ now if I g delete

• kubectl get pods -n mysql

↳ deleted by now it creates the same nothing

change
the Statefulset comes into picture

* Configmaps ↗

An API Object that lets you store configuration for other Objects to use.

- It stores configuration data (like urls, names, Environment settings) Separately from the app code

Ex. I have created Statefulset.yaml.

here I have declared configuration
dataset. ^{base} devops

If I want to change I should do it
manually.

So, to not do it manually we use configmap

Vim Configmap.yaml

kind: ConfigMap

kubectl apply -f configmap.yaml

Note: I have created Statefulset.yaml file
make it & dont change.

If any changes are required configmap
through can be done.

- Keep your app & its settings separate
- change settings without touching the app code
- Useful for non-sensitive info (secrets / password)

Kubectl get configmap -n mysql

→ to use secretfiles as binary

Secrets: →

A secret is an object that contains a small amount of sensitive data such as a password, a token, or a key.

Vim Secrets.yaml

kind: Secrets

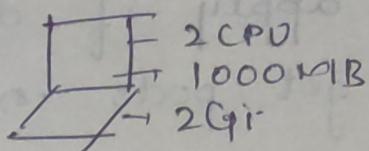
Kubectl apply -f Secrets.yaml

Kubectl get secret -n mysql

→ Scaling & Scheduling

① Resource Quotas & limits

Pod:



My laptop limit is 2CPU, 2GB, but this only 1-pod takes everything there will be nothing kept for another pool.

- In such cases we create the limits

via deployment.yaml

resources:

request:

cpu: 100m

memory: 128MiB

limits:

cpu: 200m

memory: 256MiB

↗ (request)

probes

Determine if a container is functioning correctly, whether it should be restarted or if its ready to receive traffic.

→ Types:

- liveness probe = container is live/not
- readiness " = ready to receive traffic
- startup " = is it starting up correctly

via deployment.yaml

- containerPort: 8000

- livenessProbe:

httpGet:

Same
readiness

path: /

Startup

port: 8000

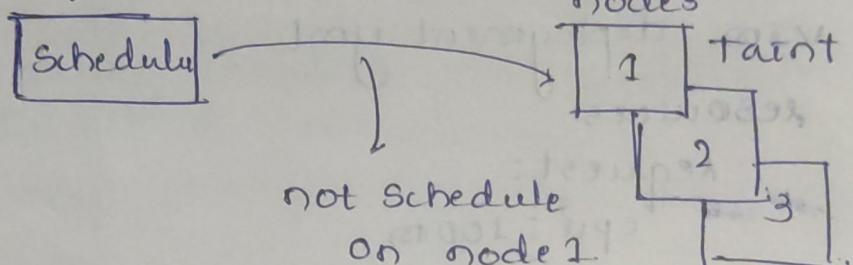
- kubectl apply -f service.yaml
- kubectl apply -f deployment.yaml
- kubectl get pods -n nginx

To check the readiness, liveness

Kubectl describe paste

Taints & Tolerations:

↳ pod not get scheduled on the node due to select nodes



• Taints are applied to nodes in a Kubernetes cluster to repel from being scheduled on those nodes.

Tolerations:

It allows a pod to run on a node with a taint.

kubernetes/ingress: ls
- pod.yaml

yaml pod.yaml

kubectl get pods -n nginx

Kubectl delete ns nginx

kubectl get nodes

↳ it displays the nodes

kubectl taint node paste pod=true:NoSchedule

-durable worker node

- `kubectl taint node two-cluster-worker2 prod=true:NoSchedule-`

- kubectl describe node -n nginx

tainted

- + how to make run tainted pods.

With pod.yaml

Spec:

tolerations:

- key: "prod"
- operator: "Equal"
- value: "true"
- effect: "NoSchedule"

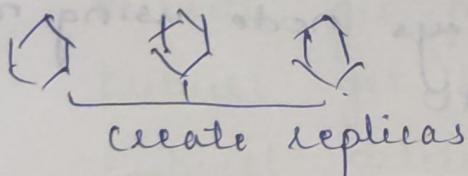
kubectl apply -f pod.yaml

" get pods -n nginx

running

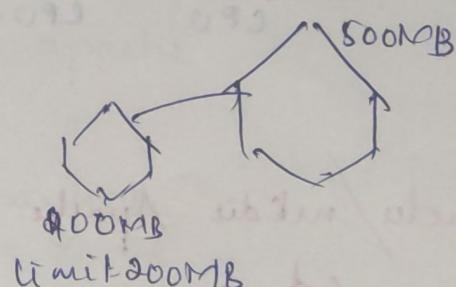
- + HPA (Horizontal pod Autoscaling) & VPA (Vertical pod Autoscaling)

HPA



- Automatically Scales the no. of pod replicas based on CPU/memory utilization

VPA



If the pod has 100 MB limits upto 200 MB if it exceeds rotate the size of the pod larger

- Automatically adjusts the CPU & memory requests / limits for containers in a pod

Metrics
CPU,
Storage
using

Open-source
component
that
allows
you to
scale
Kubernetes
workloads
based on
event-
driven
metrics

steps

- minikube addons enable metrics-server
- kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller/deployments/nginx-ingress-controller.yaml --context=minikube
↳ copy from Kube-System
↳ install metrics Server
- Edit the metrics Server deployment
kubectl -n kube-system edit deployment metrics-server
- Add the security bypass to deployment under container.args.
 - --kubelet-insecure-tls
 - --kubelet-preferred-address-types=InternalIP, Hostname, ExternalIP
- Restart the deployment
kubectl -n kube-system rollout restart deployment metrics-server

- Verify if the metrics Server is running

kubectl get pods -n kube-system

kubectl top nodes

ACTV CPU CPU % NH
S: 1002 4 4 displays node using more CPU

Kubunet/mkdir Apache

cd ..

Vim namespace.yaml

kubectl apply -f namespace.yaml

Vim deployment.yaml

Kubectl apply -f .

* To give public access
Vim service.yaml

vim hpa.yaml

apiVersion: autoscaling/v2

kind: HorizontalPodAutoscaler

metadata:

name: apache-hpa

namespace: apache

* Vertical pod Autoscaler :

Chrome = Kubernetes Autoscaler Github.

/Kubernetes/apache = git clone pasted

cd autoscaler/vertical-pod-Autosca

For installation

Github

↳ VPA

↳ Installation

Command c/p

vim vpa.yaml

kind: VerticalPodAutoscaler

apiVersion: autoscaling.k8s.io/v1

kubectl get pods -n apache

kubectl get hpa -n "

kubectl delete -f hpa.yaml

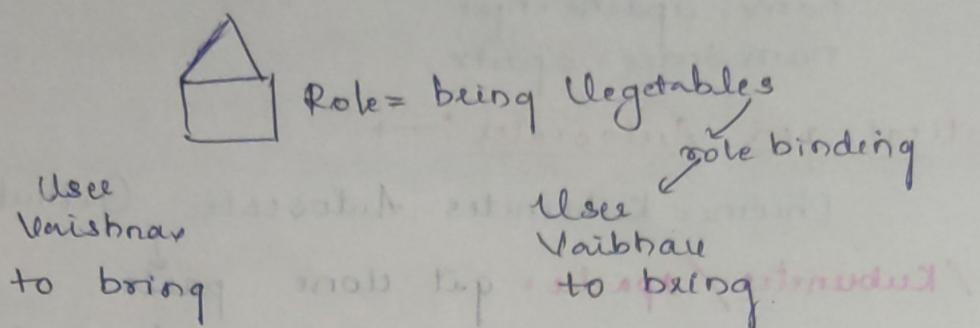
* Node Affinity :

It is a feature in Kubernetes that facilitates us to specify the rules for scheduling the pod based on the node labels.

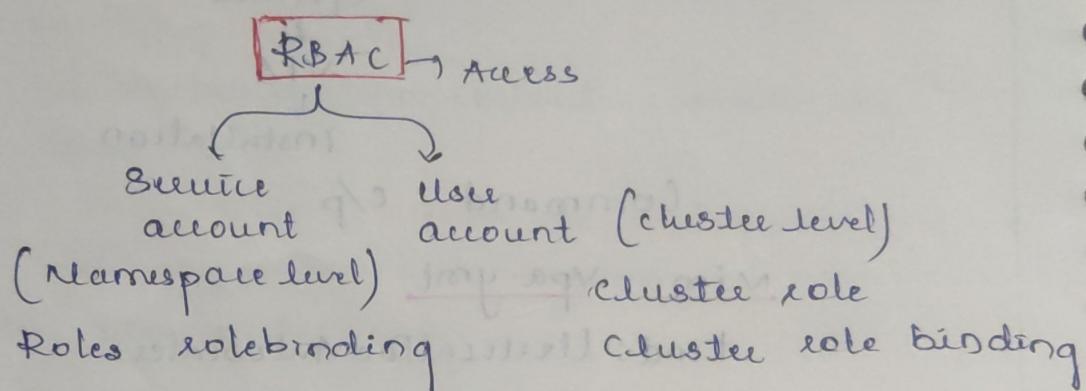
Ex:- If I'm thirsty, I need water, I atleast the water I need & the characteristics to drink water

* Cluster Administration →

* RBAC (Role Base Access Control) →



+ RBAC in Kubernetes allows you to define roles & permissions for users, groups or service accounts.



cd kubernetes

cd apache

kubectl get ns

kubectl auth whoami

" " can-i get pods

kubectl apply -f namespace.yaml

kubectl auth can-i get pods -n apache

* If I write kubectl auth delete pods -n apache, it can be deleted, the same can be done by other user. To access to the resources what I do is RBAC

```
vim role.yaml
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: apache-manage
  namespace: apache
rules:
  - apiGroups: [ "", "apps", "rbac.authorization.k8s.io" ]
    resources: ["deployment", "pod", "Service"]
    ↳ I'm giving access to these all resources.
```

Kubectl apply -f role.yaml

Kubectl get role -n apache

Vim Service-account.yaml

```
- kind: ServiceAccount
```

```
  apiVersion: v1
```

```
  metadata:
```

```
    name: apache-user
```

Kubectl apply -f service-account.yaml

Kubectl get Serviceaccount -n apache

↳ apache user
default

Kubectl auth can-i get pods -n apache

"

--as=apache-user

↳ No

Bz there is no binding

we created ["edit"]

Service role

["edit", "get", "list"]

not bind

Vim role-bind.yaml

Kind: RoleBinding

apiVersion: rbac.authorization.k8s.io/v1

- `kubectl apply -f role-bind.yaml`
- `kubectl get rolebinding -n apache`
- `kubectl auth can-i get pods -n apache --as=apache-user`
No
- `kubectl auth can-i get pods --as=apache-user -n apache`
→ Yes
- Vim role.yaml
 - ↳ make some changes = Don't allow deployment
- `kubectl apply -f role.yaml`
- `kubectl auth can-i get pods --as=apache-user -n apache`
- `kubectl auth can-i get deployment --as=apache-user -n apache`
No.

Vim role.yaml

Kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

name: apache-manager

namespace: apache

rules:

- apiGroups: ["", "apps", "rbac.authorization.k8s.io",
"batch"]

resources: ["pods", "deployments", "Services",]

verbs: ["get"]

It is the extension
of Kubernetes API

* Custom Resource Definitions (CRDs)

resources = pods, namespace, deployment

Pod → resource

apiVersion = v1

↳ version

deployment = resource

apiVersion = apps/v1

↳ group

name

* Vim CustomResourceDefinition → yaml

kind: CustomResourceDefinitions

apiVersion: apiextensions.k8s.io/v1

metadata:

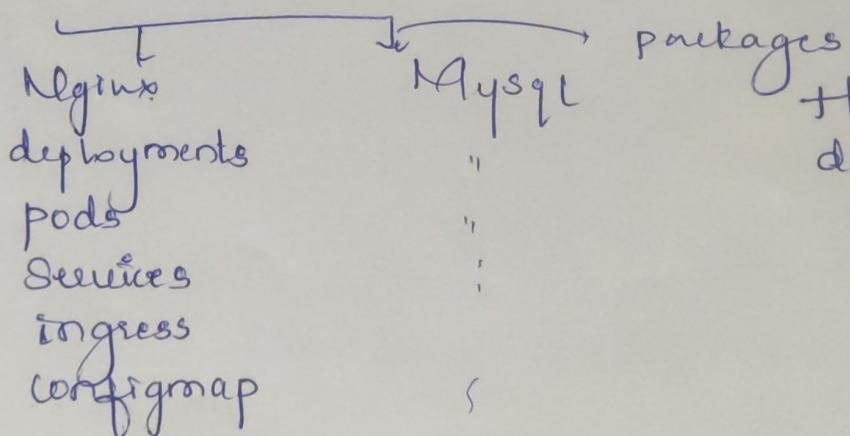
name: devops.batches

kubectl apply -f cd.yaml

* kubectl top pod -n apache

↳ displays which pod is taking more space

* Helm →



Helm = package manager for Kubernetes, with your Amazon EKS cluster to manage & deploy applications seamlessly.

Document => download

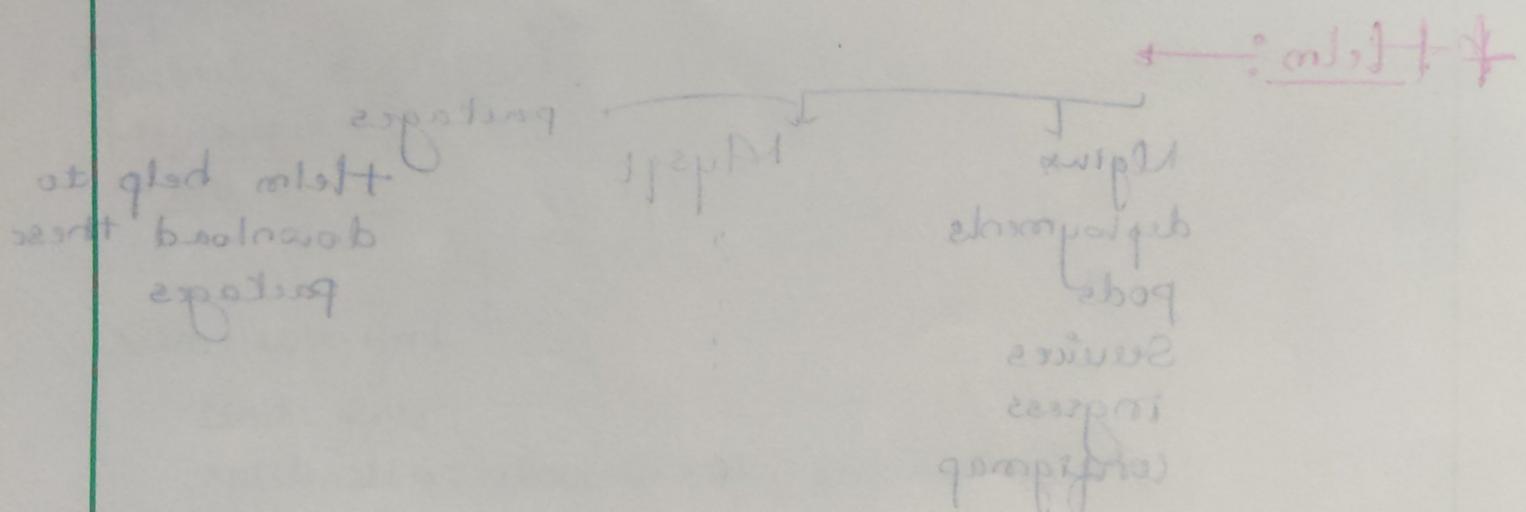
Helm is in Go language

- helm create apache-helm
 - ls
 - cd apache-helm/
 - ls
 - tree
- Li shows the tree what it has created
- ```

 chart.yaml
 charts
 templates
 - Notes.txt
 - Helpers.tpl
 - deployment.yaml
 - hpa.yaml
 - ingress.yaml
 - Service.yaml

```

single item pastet ei bog didde opgave



map items, et handlet nef nroon opgave med  
polges ð sporaar at vtekt eft aðkomu  
þleitmað er tilbúliggo

boðnaðs = tönnud  
sporaar op al ei mihi